
Gender-Bias in the Movie-Industry: Classification of Movie-Lead Genders

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 In this study, a machine learning algorithm is trained on movie-data to classify
2 the gender of the lead actor. The data-set includes number of words spoken by
3 males/females and by the lead actor, gross profit, and year of production. The
4 importance of three different features are studied to answer three main questions:
5 Do men dominate the movie industry? Has this changed over time? Do male-lead
6 movies make more money? The data included about 75 % male leads, and the
7 choice of classification method was based on both overall accuracy, and the accu-
8 racy of identifying male/female leads. The method used was Gradient Boosting,
9 with a cross-validated accuracy of around 87 % (70 % female, 94 % male). Words
10 spoken by respective genders proved to be a critical feature, while gross profit and
11 year of production made little impact. This shows of a critical gender inequality
12 in the film-industry, which is not explainable by movie profitability nor is any
13 significant improvement observed between 1970-2010.

14 1 Introduction

15 Many children, and adults, look up to their favourite movie characters. A gender imbalance in
16 movie-leads likely has some effect on how gender inequalities propagate in society. Pinpointing
17 sources of these inequalities in the movie-world could grant greater insight into what can be done to
18 minimize them. Simply noting that something is a problem is the first step to any solution. Do men
19 talk too much?

20 The article "Film Dialogue" [1] looks at the amount male vs female actors speak to detect gender
21 bias in the films. They found that male characters speak most, especially in children's movies. Is this
22 a possible source to the creation and development of gender structures in society? Has this changed
23 over the past years? Do movies where men dominate make more money? These questions are here
24 studied using various machine-learning algorithms with data-sets of parameters including words
25 spoken, production year and gross profit, and labelled after the gender of the lead-actors' gender.

26 This is an important topic to research, as it can give us insight to the gender inequalities in the film
27 industry. Given this data, we attempt to classify the gender of movie-leads and study parameter
28 importance. Notably, the data contains mainly data from men as movie-leads. As such, when
29 developing a model, there will likely be a bias towards correctly classifying men over females. While
30 historically it may be relevant to focus on correctly classifying men, as they were more frequently
31 movie-leads, to device a model reinforcing this norm might not be desirable. Rather, we believe
32 equality in accuracy has merit in itself. While fairness movie classification might not be as important
33 as for example determining whether someone should be kept in prison. Still, fairness should be

carefully considered in all machine learning projects, as the wider implications of its application can be hard to predict.

2 Methods

As mentioned in the introduction we have evaluated a few different methods for the binary classification of the gender of lead role. In the following section we explain the different methods and how they work.

(i) Logistic regression

Logistic regression [5] is a method in the family of parametric models. By using the logistic function

$$h(z) = \frac{e^z}{1 + e^z} \quad (1)$$

it is possible to define z as a linear regression model

$$z = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_p x_p = \theta^T \mathbf{x} \quad (2)$$

which squeezes the logistic function on the interval $[0, 1]$ (referred to as the *logit*). Since the model is predicting the outcome of the feature *male* or *female*, this implementation is a binary classification model. Instead of fitting to data, the sigmoid curve separates the data in the xy -plane with regards to a decision boundary. The model is trained over the data $\mathcal{T} = \{\mathbf{x}_i, y_i\}_{i=1}^n$ by numerically solving

$$\hat{\theta} = \arg \min_{\theta} \frac{1}{n} \sum_{i=1}^n \ln \left(1 + e^{-y_i \theta^T \mathbf{x}_i} \right) \quad (3)$$

The parameter vector $\hat{\theta}$ is then applied in Function 1. When evaluating test data \mathbf{x}_* , the logit outputs $\hat{y}(\mathbf{x}_*) = 1$ if $h(\mathbf{x}_*) > 0.5$ and $\hat{y}(\mathbf{x}_*) = -1$ otherwise.

The *hyperparameter* is set to $C=1$ [3] by default, but can be defined as an input for the `LogisticRegression` function to adjust for how the model should treat the weights. For large C , an obvious trade-off is overfitting since it applies large weights for the training data.

(ii) Discriminant analysis: LDA, QDA

Discriminant analysis uses the training observations to determine a boundary between response classes. The location of the boundary is determined by treating the observations of each class as samples from a multidimensional normal distribution.

Theoretically we can fit an n -dimensional normal distribution to the observations in each class. This involves calculating the mean vector and covariance matrix for each class. These determine the center and shape of the distribution.

A decision boundary is drawn where the density functions for the two distributions intersect. An implicit assumption for good performance is that the data is inherently separable; it will be hard to draw a suitable decision boundary for very mixed data.

For linear discriminant analysis (LDA) one assumes that every covariance is the same, and thus the boundary between the classes will be linear. LDA is a quick classification method, the calculations are few and simple. This function performs well when the covariance of the classes are actually the same and when the data points are not too mixed.

For quadratic discriminant analysis (QDA) one allows that different classes have different covariances. Here, the decision boundary can be quadratic, and is more flexible. QDA is still relatively quick even though it requires more memory and calculations to evaluate, store and invert the multiple covariance matrices. For QDA one can get bad estimations for the covariances if one does not have enough data, since the model will overfit to the training-data. If the data has no difference in covariance, this will perform worse than LDA. In cases with mid-sized datasets, as in this study, QDA will generally outperform LDA, due to its extra degree of adaptability.

73 **(iii) Tree-based methods: classification trees, random forests, bagging**

74 A classification tree taking in some f features and some training-data, partitions f -dimensional space
75 into disjoint parts and in each part applies a simple model (e.g. a majority vote). The partitioning is
76 often done in a way to minimize the misclassification rate in each step. The depth d of a classification
77 tree is determined by the number of splits made.

78 One way of partitioning is using the so-called Gini-index as loss function. Qualitatively, it favours
79 leaving large partitions of low quality, essentially assuming that it will be able to improve the quality
80 of the big part in future splits. This is a commonly used method, and the one used here.

81 A deeper tree generally provides low bias but high variance. A deep tree runs the risk of seriously
82 overfitting to training data. A method to reduce the variance of trees is to use a so-called ensemble
83 method. Since tree-based methods can typically be computationally cheap to find, *bagging*
84 (*bootstrapped aggregation*) of the training data \mathcal{T} can be used. Bagging is essentially drawing, with
85 replacement, data from \mathcal{T} into n new sets \mathcal{T}_i , possibly of the same size as \mathcal{T} . Now, each of these
86 bagged datasets is trained using some simple model, e.g. a classification tree, and when a decision is
87 to be made, the average output of all models is used. This way, bias is kept small, but the variance is
88 reduced.

89 One limitation of bagging with decision trees is that the bootstrapped datasets \mathcal{T}_i are correlated, so the
90 decrease in variance is limited. One potential improvement to this is the random forest (first proposed
91 by Leo Breima). The idea is essentially to perturb each tree in order to de-correlate them. This is done
92 by only considering a random subset of the input data in each split. While this (slowly) increases the
93 bias and increases the variance of each tree, the e-correlating effect is generally dominant, decreasing
94 overall variance.

95 Overall, the random forest method is easy to run in parallel, the size of bagged datasets can be tuned
96 depending on processing power, the random removal of selection points increases computation speed,
97 and the method often works well without much tuning.

98 **(iv) Boosting**

99 A similar, yet to a large part opposite approach as *bagging* is *boosting*. The same as bagging, boosting
100 is an ensemble model combining the efforts of multiple methods into a single model intended to
101 obtain more accurate predictive behaviour.

102 The idea for boosting is to use a simple model with low variance, but which is unable to map any
103 complex association between the input and output, therefore having large bias, and combine several
104 instances of the simple model, thereby reducing the bias. The combination of these weak models
105 then constitutes a single strong model, now with a lower bias, and also maintaining a low variance.
106 The concept is applicable to, and able to improve, a large number of models within machine learning.

107 The combining of models is, unlike the use of bagging, done sequentially, using the result of the
108 previous to establish the next. By re-weighting the data points depending on whether they were
109 correctly predicted or not the next model provides a result accounting to a greater degree for the
110 points misclassified by the previous. This process is repeated a predetermined number of times, each
111 model building upon the previous ones. Depending on the error of each of the models they are given
112 a coefficient determining their contribution to the one strong model.

113 Since boosting in itself is a very loose concept any choices of simple models, loss functions or
114 determination of coefficients is left unspecified by the definition, although some cases appear more
115 often. The usage of decision trees as the simple model is commonplace, and another aspect for
116 comparison with bagging and random trees, and is the simple model chosen in this instance.

117 **3 Implementations**

118 For our implementation of the methods we used sklearn `LinearDiscriminantAnalysis`
119 which implements LDA, `QuadraticDiscriminantAnalysis` which implements QDA,
120 `LogisticRegression` which implements logistic regression, `RandomForestClassifier`

121 which implements random forests, `GradientBoostingClassifier` for boosting and `kFold` which
122 splits the training data for cross validation. We chose to use the built in `sklearn` implementations
123 since it allows for simple comparison between models. Hyperparameters were initially tuned to
124 optimize for accuracy by re-running the models many times and gradually tuning parameters, and
125 later also tuned to optimize for approximately equal misclassification rates for both genders.

126 To evaluate our output we use cross validation and calculate the accuracy by calculating what
127 percentage of the data was classified correctly of all our training data. The missclassification is the
128 total number of male leads classified as females and female leads classified as male. These two
129 represent false positives and false negatives.

130 3.1 Preprocessing

131 For some methods, such as logistic regression, preprocessing of training data is needed in order to
132 perform well. There exists numerous ways of linearly transforming the training data. One method is
133 passing the data through linear combinations of functions such as \sin , \cos and \ln but in this case the
134 `StandardScaler` is used to reduce the magnitude of the dataset by letting each column sum to zero.

135 3.2 Rebalancing data

136 A method used to rebalance data was the function `SMOTE` (Synthetic Minority Over-sampling TECh-
137 nique) from the package `imblearn`. [4] As the name implies, it randomly re-samples the minority
138 class until they are of equal size. This will introduce some bias, so should be used with some caution.

139 3.3 Cross validation

140 Cross validation is a method of evaluating and comparing how well a trained classification method
141 performs, k-fold is the simplest and most used cross validation technique. It works by dividing
142 the shuffled training data set into n segments and running the learning algorithm on all data except
143 one segment and analyzing the performance (testing) on the last segment. This is then done for all
144 combinations of choosing one from n , and thus all data will be tested on once but trained on $n - 1$
145 times. Here, the `sklearn` built-in function `kFold` was used with 10 folds for our cross-validation.

146 3.4 Feature importance

147 To test feature-importance, a function to generate data-sets without all combinations of some select
148 features is used. The various models are then trained and cross-validated on all combinations. To
149 compare the models, all models were merged into one script to run sequentially.

150 4 Results

Gradient Booster Classifier Without 'Year', 'Gross'			Gradient Booster Classifier Without 'Year'			Gradient Booster Classifier Without 'Year', 'Gross'		
	female	male		female	male		female	male
female'	172	56	female'	168	48	female'	179	53
male'	82	729	male'	86	737	male'	75	732
	86.7% accuracy			87.1% accuracy			87.6% accuracy	

Figure 1: Confusion matrices from Gradient Booster Classifier. The predicted classes are denoted with an apostrophe.

151 The data in table 1 is the average of 10 times boosting for each method.




Figure 2: Accuracy and accuracy for females for 11 runs with cross validation (n=10) of an "all-male-guess" as well as 8 instances of QDA with some parameters dropped. The parameters dropped from left to right are: (2) None, (3) Year, (4) Gross, (5) Year and Gross, (6) Number words female and Number words male, (7) Year, Number words female and Number words male, (8) Gross, Number words female and Number words male, (9) Year, gross and number words female and number words male.




Figure 3: Accuracy and accuracy for females for 11 runs with cross validation (n=10) of an "all-male-guess" as well as 8 instances of GradientBoostingClassifier with some parameters dropped. The parameters dropped from left to right are: (2) None, (3) Year, (4) Gross, (5) Year and Gross, (6) Number words female and Number words male, (7) Year, Number words female and Number words male, (8) Gross, Number words female and Number words male, (9) Year, gross and number words female and number words male.

Table 1: Performance data of every method when excluding 'Gross' and 'Year'

Method	Accuracy [0,1]		
	Total	Female	Male
Log Reg	0.86910	0.61811	0.95032
QDA	0.85467	0.79528	0.87389
Tree	0.82676	0.68110	0.87389
Grad Boost	0.87680	0.70472	0.93758




Figure 4: Accuracy and accuracy for females for 11 runs with cross validation (n=10) of an "all-male-guess" as well as 8 instances of LogisticRegression with some parameters dropped. The parameters dropped from left to right are: (2) None, (3) Year, (4) Gross, (5) Year and Gross, (6) Number words female and Number words male, (7) Year, Number words female and Number words male, (8) Gross, Number words female and Number words male, (9) Year, gross and number words female and number words male.

5 Discussion

For the purpose of this discussion, the data set provided as training data is assumed to be representative of Hollywood-film as a whole, and that assumptions made about this particular subset is applicable the movie industry as a whole. Whether this is true is part of a discussion not pertinent to this project.

The reality is that the majority of men have lead roles in Hollywood movies, and this leads to a bias for our model. In our given training data of 1039 movies 785 have male lead roles. When we choose a model, we have to balance the overall performance on the data where men dominate the amount of lead roles, and the accuracy balance between classifying male vs female lead roles correctly. If we want our model to perform best overall on our given training data, we can choose a model that performs really well on classifying males and rather bad on classifying female leads which gives us a quite good model based on the fact that our training data has a majority of male leads.

However, as we do not know the distribution in the final test set, we have to make a choice when selecting a final model. We want to make sure the model performs overall well, and does not act unreasonably unfair, and therefore we have chosen to use a model that performs better on both classifying male and female leads, even if this means a slight drop in overall accuracy on classifying specifically our training set.

168 5.1 Rebalancing data

169 During testing, re-balancing of data-sets using SMOTE proved inefficient in improving our model.
170 Further tuning could likely have yielded more even performances for F/M classification, but the re-
171 sampling proved to introduce significant bias for the random forest which significantly over-performed
172 during testing, though likely just due to the heavy re-sampling.

173 5.2 Feature importance

174 For the different models it can be observed in figure 2, 3, and 4 respectively how well the model
175 performs when trained on data including and excluding combinations of three features. We see that
176 including the number of words spoken by female and male non-lead roles increases the performance
177 of our model, and thus a strong association can be derived between the gender of the lead of a movie
178 and by which genders all the non-lead words were spoken. For the other two features, both the feature
179 'Gross' and 'Year', the result closely resembles that of which when no features are excluded and is
180 therefore seen to have a very minor effect on the prediction.

181 Along with changes in society over the past decades it could be expected that gender bias in the film
182 industry would have come to see notable changes over the years. This is a correlation not found in
183 the process of this project. The simple fact that the release 'Year' feature is of little to no significance
184 when it comes to learning an accurately predicting the gender of the lead actor is largely equivalent to
185 saying that the gender bias is equal for any given year and has therefore not changed over time.

186 The income of a film is to be considered a strong incentive for the making of a movie and evaluating
187 the factors that may impact the films gross-product would be part of the decision making within the
188 film industry. It has been noted that the feature 'Gross' is largely insignificant to the prediction of
189 whether the lead actor is male or female and an assumption that female-lead films are associated with
190 a lesser box-office is as such erroneous. It should however be noted that year and gross product are
191 correlated but since the exclusion of both still yield the an unaffectedly good model they can still
192 both be considered insignificant.

193 5.3 Conclusions and wider reflections

194 Selecting a "best" method has not proven trivial. On average, the various methods show similar
195 performance: overall accuracy lying between ~83-88%. However, when considering the variance
196 of some methods (i.e. not averaging over cross-validation results), some models show outliers with
197 rather bad performance. Specifically QDA displays this behaviour. This signifies a higher variance
198 and as such a higher degree of susceptibility to unique test sets which to greater extent differ from
199 the original training set. Another critical benchmark for choosing model is accuracy per gender.
200 Most models performed significantly better on male-leads than on female ones; e.g. the chosen
201 model, compares 70/94% F/M accuracy, compared to QDA 79/90%. This, combined with a high
202 average accuracy, would be a strong argument for QDA, but the great variance displayed in figure 2
203 is deterring. Comparing to figure 3, a great reduction in outliers is noted.

204 To conclude, the study of feature importance shows that words spoken by respective genders is a
205 critical feature for determining lead gender: Males dominate movies where they lead. Gross profit
206 and year of production on the other hand made very little impact and sometimes removing them even
207 improved performance slightly. This displays that not only is there a gender equality issue in the
208 Hollywood film industry, but any notion that movies starring female lead actresses implies a lower
209 gross profit is erroneous. Also incorrect is the idea that this is an issue already being resolved; while
210 the data is not all too recent (most recent data 2010), no improvement appears detectable at all in the
211 frequency of female leads.

212 References

- 213 [1] Hanah Anderson and Matt Daniels (April 2016) *Film Dialogue* url: [https://pudding.cool/2017/03/](https://pudding.cool/2017/03/film-dialogue/)
214 [film-dialogue/](https://pudding.cool/2017/03/film-dialogue/)
- 215 [2] Ling Liu, M. Tamer Özsu (2009) *Encyclopedia of Database Systems* url: [https://link.springer.com/](https://link.springer.com/referenceworkentry/10.1007%2F978-0-387-39940-9_565)
216 [referenceworkentry/10.1007%2F978-0-387-39940-9_565](https://link.springer.com/referenceworkentry/10.1007%2F978-0-387-39940-9_565)
- 217 [3] Scikit-learn developers (BSD License) (2021) *sklearn.base: Base classes and utility functions* url: [https:](https://scikit-learn.org/stable/modules/classes.html#module-sklearn)
218 [//scikit-learn.org/stable/modules/classes.html#module-sklearn](https://scikit-learn.org/stable/modules/classes.html#module-sklearn)
- 219 [4] The imbalanced-learn developers, *SMOTE* (2021), url: [https://imbalanced-learn.org/stable/](https://imbalanced-learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html)
220 [references/generated/imblearn.over_sampling.SMOTE.html](https://imbalanced-learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html), retrieved 2021-12-06.
- 221 [5] A. Lindholm, N. Wahlström, F. Lindsten, and T.B. Schön. *Machine Learning: A First Course for Engineers*
222 *and Scientists*. Cambridge University Press, 2022.

223 Appendix

224 Code for testing the various ML-algorithms. All code is also available on GitHub: <https://github.com/benjaminverbeek/Project-Statistical-ML>
225

```
1  import pandas as pd
2  import numpy as np
3  import matplotlib
4  import matplotlib.pyplot as plt
5  import random
6
7  import sklearn.preprocessing as skl_pre
8  import sklearn.linear_model as skl_lm
9  import sklearn.discriminant_analysis as skl_da
10 import sklearn.neighbors as skl_nb
11 from sklearn.model_selection import KFold
12 from sklearn.preprocessing import StandardScaler
13 from sklearn import tree
14 from sklearn.ensemble import BaggingClassifier, RandomForestClassifier,
   ↪ GradientBoostingClassifier
15 # For balancing trainingdata. "Synthetic Minority Over-sampling TEchnique"
16 from imblearn.over_sampling import SMOTE
17
18 def crossVal(model, X, y, print_accuracy=True, run_all_models=False, dropCols=[],
   ↪ random_state=1):
19     print(f"model is {model}")
20
21     # Split index for the folds
22     kf = KFold(n_splits = 10, shuffle = True, random_state = random_state)
23     testIndicies = []
24
25     # Initiate cumulative sum variables
26     all_predictions = []
27     all_ys = []
28     data = {'Female':[0, 0],
29            'Male':[0, 0]}
30     tot_crosstab = pd.DataFrame(data, index=['Female', 'Male'])
31
32     # Iterate over all k-folds, fit model and sum to cumulative confusion matrix
33     for train_index, test_index in kf.split(X):
34         testIndicies.append(test_index)
35         X_train, X_test = X.iloc[train_index,: ], X.iloc[test_index,: ]
36         y_train, y_test = y[train_index], y[test_index]
37
38         model.fit(X_train, y_train)
39
40         predict_prob = model.predict_proba(X_test)
41
42         prediction = np.empty(len(X_test), dtype=object)
43         prediction = np.where(predict_prob[:,0]>0.5, 'Female', 'Male')
44
45         conf_mat = pd.crosstab(prediction, y_test)
46
47         tot_crosstab = tot_crosstab + conf_mat
```

```

48
49     print(tot_crosstab)
50     # Statistics:
51     acc = (tot_crosstab['Female'][0] + tot_crosstab['Male'][1]) /
        ↳ tot_crosstab.values.sum()
52     nFem = tot_crosstab['Female'].values.sum()
53     nMale = tot_crosstab['Male'].values.sum()
54     accFem = tot_crosstab['Female'][0] / nFem
55     accMale = tot_crosstab['Male'][1] / nMale
56     percentMale = nMale / (nMale + nFem)
57
58     if print_accuracy:
59         print(f'Accuracy: {acc:.5f}')
60         print(f'Accuracy Female / Male:\t {accFem:.5f} / {accMale:.5f} \t (testdata
        ↳ contains {percentMale*100:.5f} % males)')
61
62     if run_all_models:
63         results.append((acc, accFem, accMale, model, dropCols))
64
65
66 def modelDropParams(model, X, y, dropCols=[], run_all_models=False, random_state=1):
67     """Function running model dropping some X-params. With cross-validation."""
68
69     print(f"\nResults without {dropCols}")
70     X = X.copy().drop(columns=dropCols)
71
72     crossVal(model, X, y, run_all_models=run_all_models, dropCols=dropCols,
        ↳ random_state=random_state)
73
74     #print(f'Accuracy tree: \t\t {np.mean(y_predict == y_test):.2f}')
75     #allMale = y_test.copy().replace(["Female"], "Male") # make a copy with all
        ↳ Male.
76
77     print("-----")
78
79 def allCombos(lst):
80     """Takes in a list of lists and returns a list of all combinations of list
        ↳ elements."""
81     combos = []
82     for i in range(2**len(lst)):
83         a=i
84         params = []
85         for j in range(len(lst)):
86             if a%2 == 1:
87                 params += lst[j]
88             a = a//2
89         combos.append(params)
90     return combos
91
92 def rescaleDataFrame(df):
93     scaler = StandardScaler()
94     scaled_input = scaler.fit_transform(df.values)
95     scaled_df = pd.DataFrame(scaled_input, index=df.index, columns=df.columns)
96     return scaled_df

```

```

97 #####
98
99 # Read the files into data frames
100 practiseTrain = pd.read_csv("train.csv")
101 practiceTest = pd.read_csv("test.csv")
102
103 # Split data into two frames, X and y
104 X = practiseTrain.copy().drop(columns=["Lead"])      # target
105 y = practiseTrain["Lead"]
106 #sm = SMOTE(random_state=42)
107 #X, y = sm.fit_resample(X, y)
108
109 # Rescale dataframe, can be commented to test if it gives better results or not
110 X = rescaleDataFrame(X)
111
112 # Dict with models
113 models = {
114     'boosting': GradientBoostingClassifier(n_estimators=500, learning_rate=1.0,
115     ↪ min_samples_split=0.5),
116     'LDA': skl_da.LinearDiscriminantAnalysis(),
117     'QDA': skl_da.QuadraticDiscriminantAnalysis(),
118     'random-forest': RandomForestClassifier(max_depth=5, min_samples_leaf=1,
119     ↪ class_weight="balanced"),
120     'logreg': skl_lm.LogisticRegression(solver='lbfgs', C=12, random_state=0)
121 }
122
123 # 'tree': tree.DecisionTreeClassifier(max_depth=4, min_samples_leaf=1)
124 # 'boosting': GradientBoostingClassifier(n_estimators=500, learning_rate=0.4,
125 ↪ min_samples_split=0.5),
126
127 model = models['random-forest']
128 run_all_models = True
129
130 # Declare parameters to evaluate and extract all combos
131 testParams = [["Year"], ["Gross"], ["Number words female", "Number words male"]]
132 combos = allCombos(testParams)
133 print(f"Generated {len(combos)} combinations.")
134 print("Running ML-algo. for all combos.")
135
136 # TODO: possibly add output to excel for easier report-writing? Or all just take
137 ↪ their model and write.
138 # OR save results to a dict and find max accuracy.
139 # Iterate over all combos
140
141 if run_all_models:
142     results = []
143     for model in models.values():
144         print(f'----- RUNNING MODEL: {model} -----')
145         for c in combos:
146             modelDropParams(model, X, y, dropCols=c, run_all_models=True,
147             ↪ random_state=random.randint(0,42))
148
149     print('\n#### FINAL RESULTS ####')
150     print(f'Top 5 by total accuracy: \n')
151     print(*list(reversed(sorted(results)[-5:])), sep='\n') # prints line-by-line

```

```
146     print('#####')
147     print(f'Worst 5 by total accuracy: \n')
148     print(*sorted(results)[:5], sep='\n')
149 else:
150     for c in combos:
151         modelDropParams(model, X, y, dropCols=c)
```
